

AI Meets Control Theory

From Classical Control to Intelligent Autonomous Systems

**Living Technical Report — M9 — Advanced Synthesis, Adaptive Control & ICC
Leaderboard**

Generated by the project team

Lead: puma | **Code:** puma, toku | **Docs/Research:** lava, famo

September 3, 2026

Abstract

This report is regenerated at every project milestone. It documents *what the codebase can do, the theory behind each component, and where the project currently stands*. The project is an open-source laboratory for comparing classical control, optimal control, state estimation, system identification, constrained Model Predictive Control (MPC), machine learning, reinforcement learning, adaptive control, and hybrid AI + control safety shields on dynamical systems, under the core principle: *build every method from scratch, simulate it, visualise it, and compare it honestly*.

Contents

1	Introduction and Scope	2
2	Framework Architecture	2
2.1	Package Layout	3
3	Current Capabilities	3
4	Theory: Modelling and Simulation	5
5	Theory: Classical Control — PID Architecture	5
5.1	Worked Example: Stabilising an Unstable Plant (Experiment 03)	5
6	Theory: Modern Control, Estimation, and Optimal Design	7
6.1	Pole Placement and LQR	7
6.2	Worked Example: Cart-Pole Regulation (Experiment 04)	7
6.3	State Estimation: Observers, Kalman Filtering, and Duality	8
6.4	Worked Example: LQG on the Cart-Pole (Experiment 06)	9
6.5	Nonlinear Basin of Attraction: Cart-Pole LQR Robustness (Experiment 05) . . .	10
6.6	Nonlinear Control & Hybrid Swing-Up: Spong Energy Shaping (Experiment 07)	12
6.7	System Identification & Data-Driven Control (Experiment 09)	14
6.8	Constrained Model Predictive Control (Experiment 08)	15
7	Theory: Machine Learning and Learned Dynamics (Module 06)	18
7.1	Worked Example: Planning with Learned Dynamics vs. True Model (Experiment 10)	19

8 Theory: Reinforcement Learning for Dynamical Systems (Module 07)	20
8.1 Worked Example: Tabular Q-Learning vs. Classical Control on Pendulum (Experiment 11)	21
8.2 Worked Example: The RL Zoo vs. LQR on Cart-Pole Balance (Experiment 18)	21
9 Theory: AI + Control Hybrid Architectures (Module 08)	22
9.1 Worked Example: Shielded Tabular Q-Learning on Pendulum (Experiment 12)	23
10 Theory: Advanced State Estimation, Adaptive Control, and Robotics Capstones (Module 09)	24
10.1 Worked Example: Crazyflie 2.0 Flying a Lemniscate (Experiment 14)	24
10.2 Worked Example: EKF Output Feedback under Noisy Partial Sensing (Experiment 15)	25
10.3 Worked Example: EKF vs. UKF on Nonlinear Pendulum Estimation (Experiment 16)	26
10.4 Worked Example: Adaptive vs. Fixed Control on a Drifting Plant (Experiment 17)	27
11 Theory: The Intelligent Control Challenge (ICC) Benchmark (Module 10)	28
11.1 Worked Example: Multi-Paradigm Challenge Leaderboard (Experiment 19)	28
12 Project Status and Roadmap	29

1 Introduction and Scope

The central question of the project is: **when should we use classical control, when should we use AI, and when should we use both?** Rather than assuming a neural network or an RL agent is the answer, every problem is first attacked with PID, state feedback, observers, LQR, Kalman filtering (linear, EKF, UKF), system identification, adaptive control (MRAC), nonlinear energy shaping, constrained MPC, learned neural dynamics, deep reinforcement learning (DQN, PPO), and safety shields, and each method is measured against the others under identical initial conditions, sensor suites, noise profiles, disturbances, and actuator constraints.

Every topic follows the standardized engineering cycle:

THEORY $\rightarrow$ **DERIVATION** $\rightarrow$ **IMPLEMENTATION** $\rightarrow$ **SIMULATION**
 $\rightarrow$ **VISUALISATION** $\rightarrow$ **VALIDATION** $\rightarrow$ **COMPARISON** $\rightarrow$ **EXPERIMENT**

Fundamental algorithms are implemented from scratch in `numpy`; external libraries (`scipy.linalg`, `scipy.signal`, `control`, `gymnasium`, `stable-baselines3`) are used strictly for verification and benchmarking compliance.

2 Framework Architecture

The reusable library is the Python package `aimct` (`src/aimct/`). Its architecture separates distinct concerns behind minimal, strongly-typed interfaces so that dynamical systems, controllers, state estimators, neural network models, reinforcement learning environments, safety shields, and the simulator compose seamlessly.




Figure 1: Core framework architecture. The numerical simulation engine drives a controller (`update()`) and continuous plant or learned neural surrogate (`dynamics()`) via fixed-step 4th-order Runge–Kutta with Zero-Order Hold (ZOH).

2.1 Package Layout

```

src/aimct/
  systems/      DynamicalSystem base + LinearSystem, MassSpringDamper,
                Pendulum, CartPole, PlanarQuadrotor (Crazyflie 2.0), DCMotor
  simulate.py   rk4_step(), simulate() -> Trajectory(t, x, u, y)
  controllers/  PID, StateFeedback, LQR, ObserverFeedback, MRAC,
                EnergyShapingSwingUp, HybridSwingUpLQR,
                LinearMPC (with preview), SamplingMPC (CEM), _qp.solve_qp
  estimation/   LuenbergerObserver, KalmanFilter (LQE/FARE), DiscreteKalmanFilter,
                ExtendedKalmanFilter, UnscentedKalmanFilter, observability_matrix
  sysid/        least_squares_id, dmcdc, to_continuous (block logm), prediction_error
  ml/           MLP (backprop + Adam), LearnedDynamics (grey-box / residual)
  rl/           ControlEnv, Discretizer, QLearning, DQN, REINFORCE, PPO
  hybrid/       ShieldedController (switch/filter blends, predicate helpers)
  benchmarks/  metrics.py (13 metrics), harness.py, sweep.py, challenge.py (ICC
    engine)
  plot_style.py Okabe-Ito palette + publication-ready 4-panel comparison figure

```

3 Current Capabilities

Table 1 summarises the implemented components, their mathematical scope, and the corresponding verification test coverage as of M9 — Advanced Synthesis, Adaptive Control & ICC Leaderboard.

Component	What it provides	Tests
<code>aimct.systems</code>	Continuous-time ODE models with a common <code>dynamics(t, x, u)</code> interface, full-state or output projections, and analytical/numerical <code>linearize()</code> . Ships: <code>LinearSystem</code> , <code>MassSpringDamper</code> , <code>Pendulum</code> , <code>CartPole</code> , <code>PlanarQuadrotor</code> (Crazyflie 2.0), <code>DCMotor</code> .	✓
<code>aimct.simulate</code>	Fixed-step RK4 integrator with zero-order hold; <code>Trajectory</code> container; hard actuator saturation; measurement and disturbance injection hooks.	✓
<code>aimct.controllers.PID</code>	From scratch: filtered derivative-on-measurement ($D(s) = \frac{K_d s}{\tau_d s + 1}$), conditional-integration anti-windup clamping, setpoint weighting.	✓
<code>aimct.controllers.StateFeedback</code>	Arbitrary closed-loop pole placement via Ackermann's formula (SISO) with feedforward steady-state gain scaling k_r .	✓
<code>aimct.controllers.LQR</code>	Infinite-horizon optimal regulator; Continuous Algebraic Riccati Equation (CARE) solved from scratch via Hamiltonian Schur decomposition, cross-checked with <code>scipy.linalg</code> .	✓
<code>aimct.controllers.MRAC</code>	Model Reference Adaptive Control: Lyapunov-derived parameter adaptation tracking a reference model under drifting or uncertain plant parameters.	✓
<code>aimct.controllers.ObserverFeedback</code>	Unified output feedback combining any state estimator (Luenberger, linear KF, EKF, UKF) with static feedback gain $u = -K\hat{x}$.	✓
<code>aimct.controllers.SwingUp</code>	Nonlinear Spong energy shaping (<code>EnergyShapingSwingUp</code>) with partial feedback linearisation, plus finite-state hysteresis supervisor (<code>HybridSwingUpLQR</code>).	✓
<code>aimct.controllers.LinearMPC</code>	Constrained linear MPC via condensed receding-horizon QP, active-set solver (<code>_qp.solve_qp</code>), DARE terminal cost, soft state constraints, and trajectory preview.	✓
<code>aimct.controllers.SamplingMPC</code>	Sampling-based nonlinear MPC using the Cross-Entropy Method (CEM), vectorized batch trajectory rollouts, variance floor, and CARE terminal weight.	✓
<code>aimct.estimation</code>	State estimation suite: <code>observability_matrix</code> , <code>is_observable</code> , <code>place_observer</code> , <code>LuenbergerObserver</code> , <code>solve_fare</code> , <code>KalmanFilter</code> , <code>DiscreteKalmanFilter</code> , <code>ExtendedKalmanFilter</code> , and <code>UnscentedKalmanFilter</code> ($2n + 1$ scaled sigma points, Merwe parameters, zero Jacobian linearisation).	✓
<code>aimct.sysid</code>	Data-driven system identification: <code>least_squares_id</code> , <code>dmdc</code> (rank-truncated SVD), <code>to_continuous</code> (exact block matrix logarithm), <code>model_mismatch</code> .	✓
<code>aimct.ml</code>	Machine learning suite: multi-layer perceptron (MLP) with hand-written backpropagation + Adam, and residual / grey-box forward dynamics predictor (<code>LearnedDynamics</code>).	✓
<code>aimct.rl</code>	Reinforcement learning suite: <code>Gymnasium ControlEnv</code> , <code>Discretizer</code> , <code>Tabular QLearning</code> , <code>DQN</code> (experience replay + target net), <code>REINFORCE</code> , and <code>PPO</code> (clipped surrogate objective + GAE).	✓
<code>aimct.hybrid</code>	Hybrid AI + control safety shields: <code>ShieldedController</code> (switch and filter action blends, predicate helpers <code>box_predicate</code> , <code>barrier_predicate</code> , and audit logging).	✓
<code>aimct.benchmarks.challenge</code>	The Intelligent Control Challenge (ICC) evaluation engine: Tracks 1–4 (precision, agility, stacked robustness, black-box) with standardized composite scoring rubric.	✓
<code>aimct.benchmarks.metrics</code>	13 standard step-response and effort metrics: rise time t_r , settling time t_s (2%), overshoot M_p , e_{ss} , RMSE, IAE, ITAE, ISE, energy E_u , peak $ u _{\max}$, slew rate, and saturation duty cycle.	✓ 4

4 Theory: Modelling and Simulation

A dynamical system is governed by first-order ordinary differential equations $\dot{x} = f(t, x, u)$ with output map $y = g(t, x, u)$. Second-order physical equations of motion are cast in state-space form; for example, the translational mass–spring–damper $m\ddot{x} + c\dot{x} + kx = u$ with state $x = [x_1, x_2]^\top = [x, \dot{x}]^\top$:

$$\frac{d}{dt} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ -\frac{k}{m} & -\frac{c}{m} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{1}{m} \end{bmatrix} u.$$

Jacobian Linearisation. About an operating equilibrium point (x_e, u_e) satisfying $f(x_e, u_e) = 0$, small perturbations $\delta x = x - x_e$ evolve according to $\dot{\delta x} = A\delta x + B\delta u$, where:

$$A \triangleq \left. \frac{\partial f}{\partial x} \right|_{(x_e, u_e)}, \quad B \triangleq \left. \frac{\partial f}{\partial u} \right|_{(x_e, u_e)}.$$

The `DynamicalSystem` base class evaluates A, B numerically via central differencing ($\epsilon = 10^{-6}$), while concrete models implement analytical Jacobians.

Numerical Integration. The continuous dynamics are integrated over time steps $h = \Delta t$ using explicit 4th-order Runge–Kutta (RK4):

$$k_1 = f(t_k, x_k, u_k), \quad k_2 = f\left(t_k + \frac{h}{2}, x_k + \frac{h}{2}k_1, u_k\right), \quad k_3 = f\left(t_k + \frac{h}{2}, x_k + \frac{h}{2}k_2, u_k\right), \quad k_4 = f(t_k + h, x_k + hk_3, u_k),$$

$$x_{k+1} = x_k + \frac{h}{6} (k_1 + 2k_2 + 2k_3 + k_4).$$

Holding control input $u(t) \equiv u_k$ constant across each interval $[t_k, t_k + h)$ reflects the physical Zero-Order Hold (ZOH) of digital microcontrollers.

5 Theory: Classical Control — PID Architecture

The continuous-time PID control law with filtered derivative is:

$$u(t) = K_p e(t) + K_i \int_0^t e(\tau) d\tau + K_d \frac{de(t)}{dt}, \quad e(t) = r(t) - y(t).$$

In discrete time with sampling interval h , derivative action is applied directly to the measured output with low-pass filter time constant $\tau_d = \frac{K_d}{N K_p}$ ($N = 10$) to prevent high-frequency noise amplification:

$$e_k = r_k - y_k, \quad D_k = \frac{\tau_d}{\tau_d + h} D_{k-1} + \frac{K_d}{\tau_d + h} (y_k - y_{k-1}), \quad u_{\text{unsat},k} = K_p e_k + I_k - D_k.$$

Anti-Windup Protection. When actuators reach physical limits ($|u| \leq u_{\text{max}}$), *conditional integration (clamping)* freezes the integrator update ($I_{k+1} = I_k$) whenever u_{unsat} is saturated and e_k has the same sign as u_k , preventing destructive overshoot upon setpoint arrival.

5.1 Worked Example: Stabilising an Unstable Plant (Experiment 03)

We evaluate the controller on an open-loop unstable plant with a Right-Half Plane (RHP) pole:

$$\ddot{y}(t) - \omega_0^2 y(t) = u(t) + d(t), \quad \omega_0 = 2.0 \text{ rad/s}, \quad G(s) = \frac{1}{s^2 - 4}.$$

The system is tested under a unit step reference $r = 1.0 \text{ rad}$, a step input disturbance torque $d = 0.5 \text{ N} \cdot \text{m}$ applied at $t = 5.0 \text{ s}$, and actuator saturation $|u(t)| \leq 10.0 \text{ N} \cdot \text{m}$.

Controller	Rise t_r [s]	Settle t_s [s]	Overshoot M_p	Steady e_{ss}	IAE	ITAE	Energy E_u	Sat. %
P-only ($K_p = 20$)	0.275	10.0	$9.54 \times 10^9\%$	6.04×10^7	4.77×10^7	4.53×10^8	984.3	97.62%
PD ($K_d = 5.66$)	0.389	10.0	28.19%	0.2812	2.811	13.63	316.3	1.60%
PID (no AW)	0.334	6.162	53.08%	3.90×10^{-5}	0.9274	1.043	262.2	5.55%
PID + Anti-Windup	0.362	6.161	39.43%	3.90×10^{-5}	0.7979	0.876	245.1	1.61%

Table 2: Benchmark results for Experiment 03 (regenerated from experiments/03_pid_stabilizes_unstable/table.md).

Key Physical Insights:

- **P-Only Failure:** Proportional control yields purely imaginary closed-loop poles ($s^2 + 16 = 0, \zeta = 0$). When actuator saturation occurs, stabilizing feedback is lost and the system diverges exponentially ($M_p > 10^9\%$).
- **PD DC Gain Offset:** PD provides damping ($\zeta = 0.7071$), but the closed-loop DC gain is $\frac{K_p}{K_p - \omega_0^2} = \frac{20}{16} = 1.25$, creating an unavoidable **25% reference tracking error** ($y_{ss} = 1.25$ rad) before disturbance, and $e_{ss} = 0.2812$ rad under load.
- **PID Integral Action:** Integral action ($K_i = 15.0$) drives steady-state error strictly to zero ($e_{ss} \approx 0$).
- **Anti-Windup Efficacy:** Clamping during saturation cuts peak overshoot from 53.08% $\rightarrow$ 39.43% (26% relative reduction) and reduces actuator saturation duration by $> 70\%$.




Figure 2: Experiment 03 benchmark comparison. (a) Output tracking $y(t)$, (b) Control effort $u(t)$ with saturation limits $\pm 10 \text{ N} \cdot \text{m}$, (c) Tracking error $e(t)$, and (d) Phase portrait $(y, \dot{y})$.

6 Theory: Modern Control, Estimation, and Optimal Design

6.1 Pole Placement and LQR

For controllable linear systems $\dot{x} = Ax + Bu$, full-state feedback $u = -Kx + k_r r$ assigns arbitrary closed-loop poles $\sigma(A - BK)$. **StateFeedback** computes K via Ackermann's formula:

$$K = \begin{bmatrix} 0 & \dots & 0 & 1 \end{bmatrix} C(A, B)^{-1} \Delta_{\text{des}}(A), \quad k_r = -\frac{1}{C(A - BK)^{-1}B}.$$

Linear Quadratic Regulator (LQR). Rather than guessing pole locations, LQR computes the optimal feedback gain $K = R^{-1}B^\top P$ minimising the infinite-horizon quadratic cost:

$$J = \int_0^\infty \left(x(t)^\top Q x(t) + u(t)^\top R u(t) \right) dt, \quad Q \succeq 0, \quad R \succ 0.$$

The kernel $P = P^\top \succ 0$ uniquely solves the **Continuous Algebraic Riccati Equation (CARE)**:

$$A^\top P + PA - PBR^{-1}B^\top P + Q = 0.$$

We solve CARE from scratch via the stable invariant subspace of the Hamiltonian matrix:

$$\mathcal{H}_{\text{ham}} = \begin{bmatrix} A & -BR^{-1}B^\top \\ -Q & -A^\top \end{bmatrix} \in \mathbb{R}^{2n \times 2n}.$$

Continuous full-state LQR provides guaranteed robustness margins: gain margin $G_m \in [1/2, \infty)$ (-6 dB to $+\infty$ dB) and phase margin $\Phi_m \geq 60^\circ$ in all channels.

6.2 Worked Example: Cart-Pole Regulation (Experiment 04)

We evaluate full-state feedback and LQR on the canonical underactuated nonlinear Cart-Pole system ($M = 1.0$ kg, $m = 0.1$ kg, $\ell = 0.5$ m, $g = 9.81$ m/s²). Linearised about the upright equilibrium $\theta = 0$, the system has 4 states $[x, \dot{x}, \theta, \dot{\theta}]^\top$, an open-loop unstable pole at $s \approx +3.97$ rad/s, and a double integrator on cart position.

Controllers are designed on the linearised model and executed on the **true nonlinear plant** starting from a tilted initial state $\theta_0 = 0.10$ rad $\approx 5.7^\circ$ with actuator saturation $|u(t)| \leq 20.0$ N:

- **Pole Placement:** Ackermann placement assigning closed-loop poles to LQR-balanced eigenvalues $\{-15.62, -3.19, -1.31 \pm 1.08j\}$.
- **LQR (Balanced):** $Q = \text{diag}(10, 1, 100, 10)$, $R = 0.1 \implies K = [-10.000, -12.871, -87.138, -23.223]$.
- **LQR (Soft):** $Q = \text{diag}(1, 0.1, 10, 1)$, $R = 1.0 \implies K = [-1.000, -2.047, -30.846, -7.868]$.
- **PID (Angle Only):** Single-loop PID on pole angle θ ($K_p = -180, K_d = -22$). Cart position x is uncontrolled.

Controller	Settle θ t_s [s]	RMSE θ [rad]	Energy E_u [N ² s]	Peak $ u _{\text{max}}$ [N]	Peak $ x _{\text{max}}$ [m]	Status
Pole Placement	1.05	0.0171	2.34	8.71	0.162	Stable
LQR (Balanced)	1.05	0.0171	2.34	8.71	0.162	Stable
LQR (Soft)	1.70	0.0223	0.959	3.08	0.321	Stable
PID (Angle Only)	0.20	0.0110	8.76	18.00	0.823	Stable (Cart Drifts)

Table 3: Benchmark results for Experiment 04 on nonlinear Cart-Pole (from `experiments/04_lqr_vs_pole_placement_cartpole/table.md`).

Key Engineering Takeaways:

1. **Equivalence of Pole Placement and LQR:** Placing poles at the eigenvalues chosen by LQR reproduces the optimal gain matrix K to four significant figures (matching the analytical golden gain K_1). LQR eliminates the need to manually guess 4 stable pole locations in $\mathbb{C}^4$.

2. **The Q/R Effort Dial:** Softening the cost weights ($R = 1.0$ vs 0.1) trades a $1.05\text{ s} \rightarrow 1.70\text{ s}$ settle time for a $2.4\times$ **reduction in control energy** ($2.34 \rightarrow 0.96\text{ N}^2\text{s}$) and cuts peak force from $8.71\text{ N} \rightarrow 3.08\text{ N}$.
3. **Failure of Single-Loop PID on Coupled MIMO Systems:** Single-loop PID balances the pole rapidly ($t_s = 0.20\text{ s}$), but because it lacks state feedback on the cart, the cart drifts to $x = 0.823\text{ m}$ ($5\times$ LQR excursion) and demands 18.0 N peak force (near the 20 N limit). Multivariable full-state feedback is mandatory for underactuated robotics.
4. **Nonlinear Validity Domain:** For initial angle $\theta_0 = 0.1\text{ rad} \approx 5.7^\circ$, state trajectories stay strictly within the linear basin of attraction ($\approx 0.38\text{ rad}$), enabling linear LQR to stabilize the full nonlinear plant.




Figure 3: Experiment 04 benchmark comparison on nonlinear Cart-Pole. (a) Pole angle regulation $\theta(t)$, (b) Control force $u(t)$ with actuator bounds $\pm 20\text{ N}$, (c) Cart position $x(t)$ revealing uncontrolled drift under single-loop PID, and (d) Pole phase portrait $(\theta, \dot{\theta})$.

6.3 State Estimation: Observers, Kalman Filtering, and Duality

In physical applications, directly measuring the complete state vector $x(t) \in \mathbb{R}^n$ is rarely possible due to sensor cost, weight, or accessibility. The `aimct. estimation` module provides full-state reconstruction from partial, noisy outputs $y(t) = Cx(t) + v(t)$.

The Full-State Luenberger Observer. Driven by plant input $u(t)$ and measured output $y(t)$:

$$\dot{\hat{x}}(t) = A\hat{x}(t) + Bu(t) + L(y(t) - C\hat{x}(t)),$$

where $e(t) \triangleq x(t) - \hat{x}(t)$ evolves autonomously as $\dot{e}(t) = (A - LC)e(t)$. By the **Separation Principle**, the eigenvalues of the combined plant-observer system decouple into:

$$\det \begin{bmatrix} sI - (A - BK) & -BK \\ 0 & sI - (A - LC) \end{bmatrix} = \det(sI - (A - BK)) \cdot \det(sI - (A - LC)).$$

Mathematical Duality: LQR Control vs. Kalman Filtering (LQE). State estimation and optimal control are exact mathematical duals:

$$\text{LQR Regulator (CARE): } A^\top P + PA - PBR^{-1}B^\top P + Q = 0, \quad K = R^{-1}B^\top P$$

$$\text{LQE Estimator (FARE): } AP + PA^\top - PC^\top R_v^{-1}CP + Q_w = 0, \quad L = PC^\top R_v^{-1}$$

where $Q_w \succeq 0$ is process disturbance covariance and $R_v \succ 0$ is measurement noise covariance.

Discrete Kalman Filter (DKF). For sampled systems ($x_{k+1} = A_d x_k + B_d u_k + w_k$), the filter executes predict-update recursions using the numerically robust **Joseph-form** covariance update:

$$\hat{x}_{k|k} = \hat{x}_{k|k-1} + K_k(y_k - C_d \hat{x}_{k|k-1}), \quad P_{k|k} = (I - K_k C_d)P_{k|k-1}(I - K_k C_d)^\top + K_k R_v K_k^\top.$$

Sensor-Suite Observability on Cart-Pole. Evaluating the observability matrix $\mathcal{O} = [C^\top, A^\top C^\top, (A^\top)^2 C^\top, \dots]^\top$ reveals profound differences across sensor configurations:

Sensor Configuration	Measurement y	Rank($\mathcal{O}$)	Physical / Structural Mechanism
Full State	$[x, \dot{x}, \theta, \dot{\theta}]^\top$	4	Direct sensing of all coordinates.
Cart Position Only	$y = x$	4 (Observable)	Cart acceleration $\ddot{x}$ dynamically couples to pole angle θ .
Pole Angle Only	$y = \theta$	2 (Unobservable)	Rigid-body rail translation produces zero torque on the pole.
Dual Encoders	$y = [x, \theta]^\top$	4 (Observable)	Direct kinematic sensing; optimal numerical conditioning.

Table 4: Observability rank of the 4-state Cart-Pole under candidate sensor suites.

6.4 Worked Example: LQG on the Cart-Pole (Experiment 06)

Full-state LQR assumes all four states are measured with zero noise. In physical hardware, velocity sensors are noisy or omitted. Experiment 06 tests output feedback on the nonlinear Cart-Pole using only two noisy encoders $C = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$ with zero-mean Gaussian noise $\sigma_x = 5.0 \text{ mm}$ and $\sigma_\theta = 8.7 \text{ mrad} \approx 0.5^\circ$, regulating from $\theta_0 = 0.08 \text{ rad} \approx 4.6^\circ$ under actuator limits $|u(t)| \leq 20.0 \text{ N}$.

Controller	Settle θ t_s [s]	RMSE θ [rad]	Energy E_u [N ² s]	Peak $ u _{\max}$ [N]	Status
LQR (Full State, Clean)	0.942	0.0150	1.49	6.97	Ideal Baseline
LQG (Optimal Filter)	2.070	0.0249	2.21	3.28	Optimal Smooth
Luenberger + K	0.496	0.0165	16.20	15.10	Severe Noise Chattering
LQG (Overconfident $V/100$)	0.828	0.0215	2.97	4.80	Noise Amplified

Table 5: Benchmark results for Experiment 06 comparing LQG, Luenberger observer, and clean LQR under measurement noise (from `experiments/06_lqg_vs_lqr_measurement_noise/table.md`).

Key Physical and Estimation Insights:

- The Role of the Kalman Filter (Noise Rejection over Raw Bandwidth):** The Kalman filter’s objective is to prevent sensor noise from corrupting actuator commands. LQG achieves the lowest control energy (2.21 N²s) and lowest peak force (3.28 N) of any noisy-sensor controller, with a visibly smooth control trace (Figure 4b). It settles in 2.07 s (vs. 0.94 s clean LQR) precisely because trading estimator bandwidth for noise rejection is the mathematically optimal compromise.
- The Cost of Fast Pole Placement Without a Noise Model:** Fast pole placement (Luenberger + K) achieves fast settling (0.50 s) but directly amplifies raw encoder noise into violent high-frequency actuator chattering, consuming 11× **more control energy**

(16.20 N²s) and demanding 15.1 N peak force. Fast observers without noise awareness destroy actuator longevity.

3. **Significance of Covariance Tuning (V):** An overconfident Kalman filter ($V/100$) over-trusts noisy measurements, increasing control effort (2.97 N²s) and peak force (4.80 N) while increasing vulnerability to disturbance spikes.
4. **Near-Ideal Reconstruction from Partial Sensing:** Reconstructing 4 dynamic states from just 2 noisy encoders gets within $1.7\times$ the clean-LQR angle RMSE while demanding *lower* peak force (3.28 N vs. 6.97 N), confirming that the Separation Principle operates effectively on real nonlinear systems.




Figure 4: Experiment 06 benchmark comparison under encoder noise ($\sigma_x = 5$ mm, $\sigma_\theta = 0.5^\circ$). (a) Pole angle tracking $\theta(t)$, (b) Control force $u(t)$ contrasting LQG smooth action against Luenberger violent noise chattering, (c) Tracking error, and (d) Phase portrait $(\theta, \dot{\theta})$.

6.5 Nonlinear Basin of Attraction: Cart-Pole LQR Robustness (Experiment 05)

A linear controller $u = -Kx$ designed on the upright linearisation operates within a finite **basin of attraction** on the true nonlinear plant. In Experiment 05, we quantify the boundary of recoverable initial conditions $(\theta_0, \dot{\theta}_0)$ across three reference LQR tunings under hard actuator saturation $|F| \leq 20.0$ N.

Tuning	Effort R	Measured θ_0 Edge	Lyapunov Est.	Measured $\dot{\theta}_0$ Edge	Lyapunov Est.
Balanced	0.10	0.83 rad (48°)	0.80 rad (46°)	4.4 rad/s	4.0 rad/s
Aggressive	0.01	0.92 rad (53°)	0.80 rad (46°)	5.3 rad/s	4.0 rad/s
Soft	1.00	1.00 rad (57°)	0.33 rad (19°)	5.3 rad/s	1.3 rad/s

Table 6: Measured recoverable boundaries vs. theoretical Lyapunov sub-level sets on nonlinear Cart-Pole.

Key Physical and Algorithmic Findings:

1. **Replication of Literature Envelopes:** The from-scratch CARE solver precisely replicates the canonical Lyapunov envelope (0.83 rad/4.4 rad/s measured vs. 0.80 rad/4.0 rad/s predicted).
2. **Conservatism of Quadratic Lyapunov Sets:** The sub-level set $x_0^\top P x_0 \leq c^*$ provides an inner conservative bound. For the soft tuning, the true nonlinear recoverable angle is $\sim 3\times$ larger (1.00 rad vs 0.33 rad) because the quadratic form P is not aligned with the nonlinear invariant manifold.
3. **Actuator Saturation Expands the Basin:** Unconstrained linear feedback over-commands extreme forces at large angles ($\theta > 50^\circ$), leading to rapid numerical and physical divergence. Hard clamping to ± 20 N produces smoother, recoverable trajectories.
4. **Boundary for Nonlinear Control:** Beyond $\theta_0 > 57^\circ$, linear state feedback fails definitively, establishing the explicit operating domain where **nonlinear energy-based swing-up** is required.




Figure 5: Experiment 05 2D ($\theta_0 \times \dot{\theta}_0$) basin of attraction maps comparing Balanced, Aggressive, and Soft LQR tunings under ± 20 N actuator limits on nonlinear Cart-Pole.




Figure 6: Experiment 05 1D initial angle sweep from rest ($\theta_0 \in [0, 90^\circ]$) illustrating settling time t_s (left) and control energy E_u / peak force $|u|_{\max}$ (right).

6.6 Nonlinear Control & Hybrid Swing-Up: Spong Energy Shaping (Experiment 07)

To stabilize the Cart-Pole from the downward hanging equilibrium ($\theta_0 = \pi$), linear control cannot succeed because θ_0 lies far outside the linear basin of attraction (≤ 0.83 rad). Experiment 07 implements nonlinear **Spong Energy Shaping** coupled with a hybrid finite-state supervisor (**HybridSwingUpLQR**) to pump mechanical energy into the pendulum and hand off control cleanly to LQR near upright.

Energy Shaping Mechanics. The total mechanical energy of the unactuated pendulum relative to upright is:

$$E(\theta, \dot{\theta}) = \frac{1}{2} J \dot{\theta}^2 + mg\ell(\cos \theta - 1),$$

where $E(0, 0) = 0$ corresponds to the upright target and $E(\pi, 0) = -2mg\ell$ is the downward rest state. The Spong energy-pumping law applies desired cart acceleration:

$$\ddot{x}_{\text{des}} = k_E(E(\theta, \dot{\theta}) - E_{\text{target}})\dot{\theta} \cos \theta - k_p x - k_d \dot{x},$$

which is mapped into physical cart force $u(t)$ via partial feedback linearisation:

$$u = (M + m - m \cos^2 \theta) \ddot{x}_{\text{des}} - m\ell\dot{\theta}^2 \sin \theta + mg \sin \theta \cos \theta.$$

Hysteresis Supervisor Logic. To eliminate limit-cycle chattering near the transition boundary, the hybrid supervisor switches between modes with asymmetric threshold hysteresis:

Capture into LQR if: $|\text{wrap}(\theta)| \leq \theta_{\text{capture}} (0.35 \text{ rad})$ and $|\dot{\theta}| \leq \dot{\theta}_{\text{capture}} (1.5 \text{ rad/s})$

Revert to Swing-Up if: $|\text{wrap}(\theta)| \geq \theta_{\text{release}} (0.60 \text{ rad})$.

Energy Gain k_E	Capture t_{cap} [s]	Settle t_s [s]	Switches	Energy E_u [N ² s]	Peak $ u _{\text{max}}$ [N]	Cart Excursion [m]	Status
$k_E = 6$	6.496	7.452	1	52.28	12.23	1.012	Balanced
$k_E = 10$	2.852	4.058	1	87.28	20.00	1.536	Balanced
$k_E = 14$	1.686	3.344	1	107.24	19.27	1.940	Balanced

Table 7: Benchmark results for Experiment 07 swing-up from $\theta_0 = \pi$ across energy gains $k_E \in \{6, 10, 14\}$ under $|F| \leq 20$ N (from `experiments/07_cartpole_swingup_hybrid/table.md`).

Key Nonlinear Control Insights:

1. **Universal Global Convergence:** Energy shaping successfully lifts the pendulum from $E = -2mg\ell$ to the zero-energy homoclinic orbit for all gains, and LQR catches it with zero failure (Figure 7).
2. **Single-Switch Chattering-Free Handoff:** Exactly one mode transition occurs per simulation, proving the robustness of the $0.35 \rightarrow 0.60$ rad hysteresis window.
3. **Monotonic Gain Trade-Off:** Increasing k_E from $6 \rightarrow 14$ cuts capture time $\sim 4\times$ ($6.50 \text{ s} \rightarrow 1.69 \text{ s}$) but doubles control energy ($52.28 \rightarrow 107.24 \text{ N}^2\text{s}$) and expands rail excursion ($1.01 \text{ m} \rightarrow 1.94 \text{ m}$). $k_E \approx 10$ represents the optimal Pareto knee.




Figure 7: Experiment 07 trajectory traces during swing-up from hanging ($\theta_0 = \pi$). Top panels show pole angle $\theta(t)$ and pendulum total energy $E(t) \rightarrow 0$ with the exact LQR handoff instant marked. Bottom panels show control force $u(t)$ and cart position $x(t)$.




Figure 8: Experiment 07 canonical 4-panel comparison for $k_E \in \{6, 10, 14\}$ on nonlinear Cart-Pole.

6.7 System Identification & Data-Driven Control (Experiment 09)

Model-based controllers (LQR, LQG, MPC) depend directly on the mathematical plant model (A, B) . In practical applications, physical parameters are uncertain or unmodeled. The `aimct.sysid` module estimates discrete-time state-space models (A_d, B_d) directly from input-output trajectories (X_0, U, X_1) via regularized Least-Squares and Dynamic Mode Decomposition with Control (DMDc):

$$\min_{A_d, B_d} \left\| X_1 - [A_d \ B_d] \begin{bmatrix} X_0 \\ U \end{bmatrix} \right\|_F^2 + \lambda \| [A_d \ B_d] \|_F^2.$$

Continuous-time matrices (A, B) are recovered through exact Zero-Order Hold (ZOH) block-matrix logarithm inversion:

$$\begin{bmatrix} A & B \\ 0 & 0 \end{bmatrix} = \frac{1}{\Delta t} \logm \begin{bmatrix} A_d & B_d \\ 0 & I_m \end{bmatrix}.$$

Closed-Loop Identification vs. Data Horizon (Experiment 09). We fit models on noisy closed-loop Cart-Pole rollout data ($\sigma_{\text{pos}} = 3 \text{ mm}$, $\sigma_{\text{angle}} = 0.35^\circ$) across three data budgets, evaluating LQR control designed on the estimated model against the true nonlinear plant:

Data Length	Relative Error A	Relative Error K_{id}	RMSE θ [rad]	Stable Runs	Status
300 steps (1 s)	5.86×10^0	0.992	26.77	1 / 6	Severe Model Drift / Instability
1500 steps (3 s)	4.19×10^{-1}	0.527	0.0336	4 / 6	Marginal / Oscillatory
12000 steps (24 s)	2.17×10^{-1}	0.493	0.0246	6 / 6 (100%)	Consistently Stable

Table 8: Experiment 09 Monte Carlo evaluation (6 seeds) of LQR designed on identified models (from `experiments/09_control_on_identified_model/table.md`).

Key System Identification Insights:

1. **The Data Budget Threshold:** 1 s of closed-loop data yields wildly inaccurate Jacobians ($5.86 \times$ relative error in A), causing 5/6 controllers to destabilize the plant and spin the pole off-frame (Figure 9).
2. **Closed-Loop Feedback Correlation Bias:** In closed-loop data, input $u_k = -Kx_k + v_k$ is strongly correlated with state x_k , causing the least-squares estimator error to plateau around $\approx 20\%$.
3. **LQR Robustness Absorbs Identification Residuals:** Despite 20% residual parameter error, LQR designed on 24 s data achieves 100% stabilization across all seeds, demonstrating how LQR’s $[-6 \text{ dB}, +\infty \text{ dB}]$ gain margin tolerates substantial identification error. This serves as the bridge to learned neural dynamics in Module 06.




Figure 9: Experiment 09 benchmark comparison showing LQR performance designed on identified models from 300, 1500, and 12000 steps of training data versus true model baseline.

6.8 Constrained Model Predictive Control (Experiment 08)

While LQR produces optimal linear feedback, it cannot enforce state inequality constraints during trajectory planning. In robotic systems, exceeding physical travel bounds (e.g. cart running off the rail) causes mechanical damage. **Model Predictive Control (MPC)** solves an online Finite-Horizon Optimal Control Problem (FHOCPP) at every sampling instant:

$$\min_{U=\{u_0, \dots, u_{N-1}\}} \sum_{k=0}^{N-1} (x_k^\top Q x_k + u_k^\top R u_k) + x_N^\top Q_f x_N + \rho \sum_{k=1}^N (\|\max(0, x_k - x_{\max})\|^2 + \|\max(0, x_{\min} - x_k)\|^2)$$

subject to:

$$x_{k+1} = A_d x_k + B_d u_k, \quad u_{\min} \leq u_k \leq u_{\max}.$$

Condensed Quadratic Program Formulation. By expressing predicted future states $X = S_x x_0 + S_u U$ algebraically, the optimization reduces to a dense Quadratic Program in decision variables $U \in \mathbb{R}^{mN}$:

$$\min_U \frac{1}{2} U^\top H U + (M x_0 + g)^\top U \quad \text{subject to} \quad C U \leq d + S x_0, \quad u_{\min} \leq U \leq u_{\max},$$

which is solved from scratch via our primal **active-set QP solver** (`_qp.solve_qp`) with regularized Cholesky factorization and Schur complement equality-constrained subproblems.

DARE Terminal Cost, Soft Constraints & Trajectory Preview.

- **Terminal Weight $Q_f = P_{\text{DARE}}$:** Setting the terminal weight equal to the discrete infinite-horizon Riccati solution ensures that unconstrained MPC identically reproduces discrete LQR feedback to 10^{-7} precision.

- **Soft State Constraints:** Exact quadratic penalty variables maintain recursive QP feasibility even during extreme transient disturbances where hard constraints would be temporarily violated.
- **Trajectory Reference Preview:** Anticipates future curved trajectories over the horizon N , eliminating the corner-cutting lag inherent in single-setpoint planners.

Controller	Cart Peak $ x _{\max}$ [m]	Constraint Violation [m]	Settle θ t_s [s]	Energy E_u [N ² s]	Status
LQR	0.6162	0.1162 (23.2%)	2.79	28.8	Rail Violation
MPC (Unconstrained)	0.6232	0.1232 (24.6%)	2.81	27.7	Rail Violation
MPC ($x_{\text{cart}} \leq 0.5$ m)	0.5014	0.0014 (0.28%)	2.56	46.3	Strictly Compliant

Table 9: Benchmark results for Experiment 08 on nonlinear Cart-Pole with $|F| \leq 20$ N and rail limit $|x| \leq 0.5$ m from $\theta_0 = 0.35$ rad (from experiments/08_mpc_vs_lqr_constrained_cartpole/table.md).

Key Constrained Control Takeaways:

1. **LQR Has No Mechanism for State Boundaries:** LQR and unconstrained MPC swing the cart to $x = 0.62$ m, overshooting the physical track limit by 23% (Figure 10).
2. **MPC Enforces Safety Envelopes:** Constrained MPC holds the cart strictly at the boundary ($x = 0.5014$ m), actively modifying control actions to respect safety limits.
3. **Performance Gain via Rail Exploitation:** Because Constrained MPC purposefully rides the edge of allowable travel, it brings the pole upright **0.23 s sooner** (2.56 s vs. 2.79 s) at the expense of 60% more control energy (46.3 N²s vs. 28.8 N²s).
4. **Unconstrained MPC Identically Recovers LQR:** Under zero state constraints, the first move of LinearMPC matches discrete LQR to $< 10^{-7}$, confirming algebraic consistency.




Figure 10: Experiment 08 trajectory traces highlighting Cart Position $x(t)$ with the ± 0.5 m track limit, Pole Angle $\theta(t)$, and Actuator Force $u(t)$ with the ± 20 N saturation bounds.




Figure 11: Experiment 08 canonical 4-panel comparison: Constrained MPC vs. Unconstrained MPC vs. LQR on nonlinear Cart-Pole.

7 Theory: Machine Learning and Learned Dynamics (Module 06)

When analytic equations of motion cannot be derived from first principles, machine learning allows parameterizing and learning nonlinear dynamics directly from rollout data. The `aimct.ml` module implements neural network modeling and sampling-based planning from scratch in `numpy`.

From-Scratch Multi-Layer Perceptron (MLP) & Backpropagation. The forward pass through an L -layer network is:

$$z^{(l)} = W^{(l)}a^{(l-1)} + b^{(l)}, \quad a^{(l)} = \sigma(z^{(l)}), \quad l = 1, \dots, L.$$

Analytical gradients are derived via the vector chain rule and optimized using the `**Adam**` optimizer with first and second moment bias corrections:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t, \quad v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2, \quad \hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}, \quad \theta_t = \theta_{t-1} - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t.$$

Gradients are strictly verified against central finite differences ($\|g_{\text{analytic}} - g_{\text{num}}\|_{\infty} < 10^{-7}$).

Residual & Grey-Box Dynamics Parameterization. Rather than predicting the full next state x_{k+1} directly, `LearnedDynamics` supports grey-box residual dynamics over a known baseline physics step $f_{\text{nominal}}(x, u)$:

$$x_{k+1} = f_{\text{nominal}}(x_k, u_k) + \sigma_y \odot f_{\text{NN}} \left(\frac{x_k - \mu_x}{\sigma_x}, \frac{u_k - \mu_u}{\sigma_u} \right) + \mu_y.$$

This formulation guarantees identity pass-through when network weights are small and isolates learning to unmodeled physical residuals.

Sampling-Based MPC: The Cross-Entropy Method (CEM). Given any forward model $f(x, u)$ (analytical or neural), **SamplingMPC** optimizes action sequences $U \in \mathbb{R}^{H \times m}$ over prediction horizon H via iterative Gaussian refinement:

1. Sample N_{samples} action sequences $U^{(i)} \sim \mathcal{N}(\mu, \Sigma)$.
2. Roll out each sequence in parallel through $f(x, u)$ to compute total cost $J(U^{(i)}) = \sum_{k=0}^{H-1} \ell(x_k, u_k) + x_H^\top P_{\text{CARE}} x_H$.
3. Identify the top elite fraction (10%) and update (μ, Σ) with smoothing factor α :

$$\mu \leftarrow \alpha\mu + (1 - \alpha)\mu_{\text{elite}}, \quad \Sigma \leftarrow \alpha\Sigma + (1 - \alpha)\Sigma_{\text{elite}} + \sigma_{\min}^2 I.$$

4. Apply first action $u_0 = \mu_0$ and warm-start next step by shifting μ .

The Crucial Role of the CARE Terminal Cost. Short-horizon sampling planners ($H = 30 \implies 0.6$ s) cannot anticipate open-loop instability occurring beyond the horizon. Adding the infinite-horizon Riccati kernel $x_H^\top P_{\text{CARE}} x_H$ —fitted from least-squares linearization on the same data—acts as an explicit terminal Lyapunov control surrogate, ensuring asymptotic stability without requiring long prediction horizons.

7.1 Worked Example: Planning with Learned Dynamics vs. True Model (Experiment 10)

We evaluate **SamplingMPC** (CEM) driving the nonlinear Cart-Pole from $\theta_0 = 0.20$ rad $\approx 11.5^\circ$ using a 4,804-parameter residual MLP ($[5 \rightarrow 64 \rightarrow 64 \rightarrow 4]$) trained on 80 s (4,000 steps) of closed-loop PRBS data (train MSE 1.96×10^{-4} , 1-step error 4.45×10^{-4} , 30-step error 2.50×10^{-2}), compared against planning on the true equations of motion and analytic LQR:

Controller	Settle θ t_s [s]	RMSE θ [rad]	Energy E_u [N ² s]	Peak $ u _{\max}$ [N]	Status
LQR (True Model)	1.26	0.0374	8.42	17.40	Analytic Optimum
SamplingMPC (True Model)	4.44	0.0448	10.80	6.59	Stable (CEM True)
SamplingMPC (Learned Model)	3.56	0.0462	10.20	6.00	Stable (CEM Learned)

Table 10: Benchmark results for Experiment 10 comparing CEM-MPC on learned neural dynamics vs. true model and LQR (from `experiments/10_planning_learned_vs_true_model/table.md`).

Key Machine Learning Insights:

1. **Learned-Model vs. True-Model Parity:** CEM-MPC on the learned neural network achieves performance within 3% RMSE (0.0462 vs. 0.0448) and 6% control energy of planning on the exact physical equations. Trajectories are visually indistinguishable (Figure 12).
2. **Zero Analytic Physics in the Learned Pipeline:** Data $\rightarrow$ MLP one-step model $\rightarrow$ data-fitted CARE terminal kernel $\rightarrow$ CEM plan $\rightarrow$ stable balance. No symbolic equations of motion are used anywhere.
3. **Planner Gap vs. Model Gap:** The performance difference relative to LQR (3.56 s vs. 1.26 s) is due to the stochastic nature of sampling MPC (averaging over elite distributions rather than executing a sharp bang-bang command), not errors in the learned neural dynamics.
4. **Necessity of the Terminal Kernel:** Without $x_N^\top P x_N$, the 0.6 s prediction horizon fails to stabilize the unstable cart-pole (yielding RMSE > 1.0).




Figure 12: Experiment 10 benchmark comparison: SamplingMPC on Learned Neural Model vs. SamplingMPC on True Model vs. LQR on nonlinear Cart-Pole.

8 Theory: Reinforcement Learning for Dynamical Systems (Module 07)

Reinforcement learning approaches control as optimal decision-making in an unknown Markov Decision Process (MDP) $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$. The `aimct.rl` module bridges the dynamical systems library with the modern Gymnasium API and provides complete tabular and deep RL algorithms from scratch.

Continuous Control Environment (ControlEnv). Wrapping continuous-time dynamical systems $\dot{x} = f(t, x, u)$ into a standard Gymnasium environment:

- **Observation Space:** State vector $x \in \mathbb{R}^n$ or trigonometric transforms $[\cos \theta, \sin \theta, \dot{\theta}]^\top \in \mathbb{R}^3$.
- **Action Space:** Continuous torque/force box bounds $u \in [u_{\min}, u_{\max}]$.
- **Quadratic Stage Reward:** $r(x_k, u_k) = -(x_k^\top Q x_k + u_k^\top R u_k)$.
- **Integration:** Fixed-step Runge-Kutta 4th-order (`rk4_step`).

Deep Q-Networks (DQN) & Proximal Policy Optimization (PPO).

- **DQN (`aimct.rl.DQN`):** Multi-head value approximator with experience replay buffer $\mathcal{D}$ and polyak/periodic target networks Q_{target} to stabilize Temporal Difference updates:

$$\mathcal{L}(\theta) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{D}} \left[\left(r + \gamma \max_{a'} Q_{\text{target}}(s', a'; \theta^-) - Q(s, a; \theta) \right)^2 \right].$$

- **PPO (`aimct.rl.PPO`):** Actor-Critic policy gradient with Generalized Advantage Estima-

tion (GAE) and clipped probability ratio objective:

$$L^{\text{CLIP}}(\theta) = \hat{\mathbb{E}}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right], \quad r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}.$$

8.1 Worked Example: Tabular Q-Learning vs. Classical Control on Pendulum (Experiment 11)

We evaluate tabular Q-learning against classical energy shaping and LQR on the nonlinear Inverted Pendulum ($m = 1.0 \text{ kg}$, $L = 1.0 \text{ m}$, $b = 0.1 \text{ Nms/rad}$, $g = 9.81 \text{ m/s}^2$) under torque limit $|u| \leq 4.0 \text{ N} \cdot \text{m}$ across swing-up and balancing:

Controller	Task	Min Err [rad]	Held Upright?	t_{upright} [s]	Final Err [rad]	Energy E_u	Train Samples	Parameters
Energy-Shaping + LQR	Swing-Up	0.00	Yes	4.00	0.000	45.9	0	3
Tabular Q-Learning	Swing-Up	0.02	No	9.10	1.479	101.2	450,000	61,875
LQR	Balance	0.00	Yes	0.85	0.000	11.7	0	2
Tabular Q-Learning	Balance	0.05	No	3.45	1.527	40.4	180,000	51,975

Table 11: Benchmark results for Experiment 11 comparing Tabular Q-Learning against Classical Energy-Shaping and LQR on the Inverted Pendulum (from experiments/11_qlearning_vs_classical/table.md).




Figure 13: Experiment 11 benchmark comparison: Tabular Q-Learning vs. Classical Control on Inverted Pendulum.

8.2 Worked Example: The RL Zoo vs. LQR on Cart-Pole Balance (Experiment 18)

Experiment 18 pits the complete RL suite (from tabular to modern policy gradients and external baselines) against linear LQR on Cart-Pole balance from $\theta_0 = 0.1 \text{ rad}$:

Agent / Controller	Greedy Return	Held Steps (/200)	Env Steps	Train Time [s]	Parameters
LQR (Analytic CARE)	-0.3	200.0	0	0.0	4
Tabular Q-Learning	-55.5	92.0	300,000	7.1	354,375
DQN (From Scratch)	-2.6	200.0	90,177	68.7	4,805
PPO (From Scratch)	-0.3	200.0	240,000	72.6	4,545
PPO (Stable-Baselines3)	-44.9	52.5	120,000	129.6	9,091

Table 12: Benchmark results for Experiment 18 comparing the RL zoo against LQR on Cart-Pole balance (from `experiments/18_rl_zoo_vs_lqr/table.md`).

Key RL Zoo Insights:

1. **From-Scratch PPO Matches LQR Performance Exactly:** Our from-scratch PPO achieves identical return (-0.3) and holds upright for all 200 steps, proving the mathematical fidelity of our custom actor-critic implementation.
2. **Sample Efficiency vs. Analytic Design:** While PPO successfully converges to the optimal gain surface, it requires **240,000 environment steps** (≈ 1.3 hours of physical interaction) compared to **0 steps** and 0.001 s for algebraic LQR.
3. **Curse of Discretization in Tabular Q:** Tabular methods blow up to $> 350,000$ table entries and fail to balance (92/200 held) due to grid coarseness near the origin, while parametric neural networks (DQN, PPO) generalize smoothly.




Figure 14: Experiment 18 benchmark comparison: RL Zoo vs. LQR on Cart-Pole balance illustrating training reward curves, pole balance angle, and control force.

9 Theory: AI + Control Hybrid Architectures (Module 08)

The fundamental dilemma between learning-based methods and classical control is clear:

- **Learning policies (RL, learned MPC):** Capable of discovering global, non-convex maneuvers from data without manual mathematical derivations, but offer zero stability, safety, or constraint guarantees.
- **Classical controllers (LQR, Lyapunov, Energy-Shaping):** Deliver rigorous stability guarantees, verified invariant domains, and provable safety, but operate within bounded local basins and require explicit physical models.

The `aimct.hybrid` module implements **Safety Shielding** to synthesize both paradigms into certifiable hybrid systems.

The Shielded Controller Architecture (ShieldedController). Let $u_{\text{base}}(x)$ denote an uncertified primary policy (e.g. a deep RL policy or learned neural planner) and $u_{\text{fallback}}(x)$ a

certified classical backup controller. The safety shield switches control authority dynamically based on a state safety predicate $\mathcal{P}_{\text{safe}} : \mathcal{X} \rightarrow \{\text{True}, \text{False}\}$:

$$u(t) = \begin{cases} u_{\text{base}}(x(t)) & \text{if } \mathcal{P}_{\text{safe}}(x(t)) = \text{True}, \\ u_{\text{fallback}}(x(t)) & \text{if } \mathcal{P}_{\text{safe}}(x(t)) = \text{False}. \end{cases}$$

9.1 Worked Example: Shielded Tabular Q-Learning on Pendulum (Experiment 12)

In Experiment 12, we wrap the raw Q-agent from Experiment 11 in `ShieldedController`:

Controller	Min Err [rad]	Final Err [rad]	Held Upright?	t_{upright} [s]	Energy E_u	Shield Active
Tabular Q-Learning (Raw)	0.028	1.406	No	10.30	105.4	0.0%
Shielded (RL + Classical Finisher)	0.000	0.000	Yes	9.30	68.0	43.7%

Table 13: Benchmark results for Experiment 12 comparing Raw Tabular Q-Learning vs. Shielded Hybrid Control on the Inverted Pendulum (from experiments/12_shielded_qlearning/table.md).




Figure 15: Experiment 12 benchmark comparison: Raw Tabular Q-Learning vs. Shielded Hybrid Control on Inverted Pendulum.

10 Theory: Advanced State Estimation, Adaptive Control, and Robotics Capstones (Module 09)

Planar Quadrotor Equations of Motion (Crazyflie 2.0). With state $x = [x, z, \theta, \dot{x}, \dot{z}, \dot{\theta}]^\top \in \mathbb{R}^6$ and motor thrusts $u = [u_1, u_2]^\top$:

$$\ddot{x} = -\frac{u_1 + u_2}{m} \sin \theta, \quad \ddot{z} = \frac{u_1 + u_2}{m} \cos \theta - g, \quad \ddot{\theta} = \frac{(u_2 - u_1)\ell}{I_{yy}}.$$

Differential Flatness & Feedforward Generation. Flat output $\sigma = [x, z]^\top$ yields closed-form pitch and thrust feedforward:

$$\theta_r(t) = -\arctan\left(\frac{\ddot{x}_r(t)}{g + \ddot{z}_r(t)}\right) \approx -\frac{\ddot{x}_r(t)}{g}, \quad T_r(t) = m\sqrt{\ddot{x}_r(t)^2 + (g + \ddot{z}_r(t))^2}, \quad \Delta T_r(t) = \frac{I_{yy}\ddot{\theta}_r(t)}{\ell}.$$

Model Reference Adaptive Control (MRAC). When plant parameters θ_{plant} drift over time, `aimct.controllers.MRAC` adjusts control gains $\hat{K}_x(t), \hat{K}_r(t)$ dynamically so the plant tracks a stable reference model $\dot{x}_m = A_m x_m + B_m r$. The tracking error $e(t) = x(t) - x_m(t)$ drives adaptation via the Lyapunov update law:

$$\dot{\hat{K}}_x(t) = -\Gamma_x B_m^\top P e(t) x(t)^\top, \quad \dot{\hat{K}}_r(t) = -\Gamma_r B_m^\top P e(t) r(t)^\top,$$

where $P = P^\top \succ 0$ satisfies the reference Lyapunov equation $A_m^\top P + P A_m = -Q_m$.

10.1 Worked Example: Crazyflie 2.0 Flying a Lemniscate (Experiment 14)

We fly the full nonlinear Crazyflie quadrotor around an aggressive figure-8 path:

$$x_r(t) = 0.6 \sin(\omega t), \quad z_r(t) = 1.0 + 0.35 \sin(2\omega t), \quad T_{\text{period}} = 4.0 \text{ s},$$

subjected to a lateral wind gust of 30 mN (11% weight) during $t \in [5, 8] \text{ s}$.

Controller	RMS Pos Err [mm]	Max Pos Err [mm]	RMS Pitch [°]	Control Energy $\int \ \Delta u\ ^2$	Thrust Saturation
LQR + Flatness Feedforward	43.5	112.0	4.3°	0.033	0.7%
LQR Feedback Only	51.0	109.0	4.4°	0.038	0.7%
Linear MPC (Single Setpoint)	142.0	208.0	4.3°	0.026	0.1%
Linear MPC (Reference Preview)	47.9	134.0	4.5°	0.033	0.6%

Table 14: Benchmark results for Experiment 14: Crazyflie 2.0 figure-8 tracking under 11% wind gust (from `experiments/14_quadrotor_figure8_tracking/table.md`).




Figure 16: Experiment 14 benchmark comparison: Crazyflie 2.0 planar quadrotor flying a lemniscate figure-8 through a sustained wind gust.

10.2 Worked Example: EKF Output Feedback under Noisy Partial Sensing (Experiment 15)

Controller	RMS Pos Err [mm]	Max Pos Err [mm]	RMS Pitch [°]	Control Energy $\int \ \Delta u\ ^2$	RMS Vel Err [m/s]
LQR (Clean Full State)	9.04	49.0	1.26°	0.0024	—
LQR + EKF (Nonlinear Observer)	9.51	50.6	1.25°	0.0026	0.008
LQR + Finite-Difference Vel	18.32	47.4	1.89°	0.3580 (150×)	0.332 (40×)

Table 15: Benchmark results for Experiment 15: Crazyflie 2.0 output feedback tracking with EKF vs. naive differencing (from experiments/15_quadrotor_ekf_output_feedback/table.md).




Figure 17: Experiment 15 benchmark comparison: EKF Output-Feedback Quadrotor Tracking.

10.3 Worked Example: EKF vs. UKF on Nonlinear Pendulum Estimation (Experiment 16)

Regime	Filter	RMS Error [rad]	Final Error [rad]	Recovered?
Hard (sin/cos, π -off init)	EKF	6.514	6.282 ($\approx 2\pi$)	False (Trapped)
Hard (sin/cos, π -off init)	UKF	1.382	0.068	True (Converged)
Mild (sin/rate, good init)	EKF	0.021	0.026	True
Mild (sin/rate, good init)	UKF	0.021	0.026	True

Table 16: Benchmark results for Experiment 16: EKF vs. UKF on nonlinear Pendulum estimation (from `experiments/16_ekf_vs_ukf/table.md`).




Figure 18: Experiment 16 benchmark comparison: EKF vs. UKF on nonlinear Pendulum estimation.

10.4 Worked Example: Adaptive vs. Fixed Control on a Drifting Plant (Experiment 17)

We evaluate Model Reference Adaptive Control (MRAC) against nominal LQR, worst-case robust LQR, and Gain-Scheduled LQR on a mass–spring–damper where the spring constant drifts continuously by 500% ($k(t) = 1.0 \rightarrow 5.0 \text{ N/m}$ over 20 s):

Controller	RMS Error (Early $t \leq 5\text{s}$)	RMS Error (Late $t \geq 15\text{s}$)	Final Settle Err [m]	Control Energy E_u
LQR (Nominal $k = 1$)	0.2573	0.4013	0.4217	180.7
LQR (Worst-Case $k = 5$)	0.3577	0.5194	0.5406	123.3
MRAC (Adaptive)	0.0008	0.0008	0.0008	413.5
GainScheduled LQR	0.2916	0.5075	0.5396	142.3

Table 17: Benchmark results for Experiment 17: Adaptive MRAC vs. Fixed LQR on a drifting plant (from experiments/17_adaptive_vs_fixed_changing_plant/table.md).

Key Adaptive Control Takeaways:

- Asymptotic Tracking Under Massive Parameter Drift:** MRAC maintains $< 1 \text{ mm}$ tracking error (0.0008 m) across the entire trajectory despite a $5\times$ parameter shift, while static LQR incurs severe steady-state droop (0.42 m).
- The Energy Price of High-Gain Adaptation:** MRAC consumes $2.3\times$ more control energy (413.5 vs 180.7 N^2s) to continuously cancel parameter estimation error.




Figure 19: Experiment 17 benchmark comparison: Adaptive MRAC vs. Fixed LQR on drifting spring plant.

11 Theory: The Intelligent Control Challenge (ICC) Benchmark (Module 10)

The **Intelligent Control Challenge (ICC)** establishes a blind black-box benchmark for comparing paradigms under identical information constraints. A submission receives only the public problem specification (dimensions, limits, time horizon, target state) without seeing the internal plant equations, and is evaluated on nominal performance, parameter sensitivity sweeps ($\pm 30\%$), and stacked disturbance shocks.

The Composite Scoring Rubric. Performance is measured relative to a tuned reference baseline (baseline score ≈ 37):

$$\text{Score} = 100 \cdot \exp \left(- \left[0.5 \frac{\text{ITAE}}{\text{ITAE}_{\text{base}}} + 0.3 \frac{E_u}{E_{u,\text{base}}} + 0.2 \frac{\text{Slew}}{\text{Slew}_{\text{base}}} \right] \right) \cdot S_{\text{robust}},$$

where $S_{\text{robust}} \in [0.20, 1.00]$ penalizes cost blowup under parameter perturbations.

11.1 Worked Example: Multi-Paradigm Challenge Leaderboard (Experiment 19)

Five paradigm representatives are evaluated across five ICC tracks:

Submission Paradigm	Track 1 (MSD)	Track 1 (DC Motor)	Track 2 (Pendulum)	Track 2 (Cart-Pole)	Track 3 (Robustness)
Linear MPC	21.1 (PASS)	41.3 (PASS)	0.0 (FAILED)	0.0 (FAILED)	0.0 (FAILED)
LQR	7.4 (PASS)	36.8 (PASS)	0.0 (FAILED)	0.0 (FAILED)	0.0 (FAILED)
Energy+LQR Hybrid	1.8 (PASS)	0.0 (FAILED)	23.8 (PASS)	0.0 (DQ)	29.9 (PASS)
PID	0.6 (PASS)	0.0 (FAILED)	0.0 (FAILED)	0.0 (FAILED)	0.0 (FAILED)
Tabular Q-Learning	0.0 (FAILED)	0.0 (FAILED)	0.0 (FAILED)	0.0 (FAILED)	0.0 (FAILED)

Table 18: Experiment 19 Intelligent Control Challenge Leaderboard: Composite Score / 100 (from experiments/19_icc_leaderboard/leaderboard.md).

Key Leaderboard Verdicts:

1. **Optimal Control Dominates Precision:** Linear MPC achieves the highest score on DC Motor (41.3, beating the baseline) and passes Mass-Spring-Damper (21.1) by optimally scheduling control authority against actuator voltage limits.
2. **Nonlinear Energy Shaping Dominates Agility & Robustness:** The Energy+LQR hybrid is the only submission to solve Pendulum swing-up (23.8) and withstand Track 3 stacked stress (29.9).
3. **No Single Paradigm Wins Everywhere:** Precision tracking rewards constrained optimization (MPC), whereas large-basin maneuvers reward global nonlinear energy methods, highlighting the necessity of hybrid architectures.




Figure 20: Experiment 19 Intelligent Control Challenge Leaderboard: Composite score comparison across all five paradigms and benchmark tracks.

12 Project Status and Roadmap

Milestone M9 — Advanced Synthesis, Adaptive Control & ICC Leaderboard — Completed. All curriculum modules, capstone robotics showcases, deep RL agents, adaptive controllers, and benchmark challenge suites are fully implemented from scratch, tested, and validated across 18 experiments:

1. **Module 01: Mathematical Foundations** — Linear algebra, matrix exponential, Runge-Kutta integration, numerical optimization.
2. **Module 02: Dynamic System Modeling** — First-principles modeling, state-space representations, Jacobian linearisation, discretization.
3. **Module 03: Classical Control** — Filtered derivative PID, conditional anti-windup clamping, frequency response, stability margins (Experiment 03).

4. **Module 04: Modern Control & Estimation** — Controllability, observability, Ackermann pole placement, full-state Luenberger observers, Filter ARE duality, continuous/discrete Kalman filtering with Joseph update, Extended Kalman Filter (EKF), Unscented Kalman Filter (UKF) (Experiments 04, 05, 06, 15, 16).
 5. **Module 05: Optimal & Constrained Control** — Continuous Algebraic Riccati Equation (CARE) Hamiltonian solver, LQR robustness margins, constrained Model Predictive Control with active-set QP solver and reference preview (Experiments 08, 14).
 6. **Module 06: System ID & Machine Learning** — Least-squares state-space ID, DMDc, matrix logarithm ZOH inversion, from-scratch MLP with Adam, residual/grey-box LearnedDynamics, SamplingMPC (CEM) with CARE terminal cost (Experiments 07, 09, 10).
 7. **Module 07: Reinforcement Learning** — Gymnasium ControlEnv adapter, state/action discretization, Tabular Q-Learning, DQN, REINFORCE, and PPO (Experiments 11, 18).
 8. **Module 08: AI + Control (Hybrid Architectures)** — Supervisory safety shielding, action filtering, barrier predicates, auditable intervention logging (Experiment 12).
 9. **Module 09: Advanced State Estimation, Adaptive Control & Robotics** — Full 6-state Planar Quadrotor (Crazyflie 2.0), differential flatness feedforward trajectory generation, Bryson-rule scaling, EKF/UKF output feedback, and Model Reference Adaptive Control (MRAC) (Experiments 14, 15, 16, 17).
 10. **Module 10: The Intelligent Control Challenge** — Standardized multi-track benchmarking engine, robustness sweeps, and cross-paradigm leaderboard (Experiment 19).
- CI passes on Python 3.10–3.13 with **305 passing unit tests**.

Reproduction Commands.

```
git clone https://github.com/zalithomas-ui/ai-meets-control-theory.git
cd ai-meets-control-theory
pip install -e ".[dev,ml,xcheck]"
python -m pytest -q
python experiments/03_pid_stabilizes_unstable/run.py
python experiments/04_lqr_vs_pole_placement_cartpole/run.py
python experiments/05_cartpole_basin_of_attraction/run.py
python experiments/06_lqg_vs_lqr_measurement_noise/run.py
python experiments/07_cartpole_swingup_hybrid/run.py
python experiments/08_mpc_vs_lqr_constrained_cartpole/run.py
python experiments/09_control_on_identified_model/run.py
python experiments/10_planning_learned_vs_true_model/run.py
python experiments/11_qlearning_vs_classical/run.py
python experiments/12_shielded_qlearning/run.py
python experiments/14_quadrotor_figure8_tracking/run.py
python experiments/15_quadrotor_ekf_output_feedback/run.py
python experiments/16_ekf_vs_ukf/run.py
python experiments/17_adaptive_vs_fixed_changing_plant/run.py
python experiments/18_rl_zoo_vs_lqr/run.py
python experiments/19_icc_leaderboard/run.py
```